

Time Series Analysis in Python | Time Series Forecasting | Data Science with Python | Edureka

eureka! channel

<https://www.youtube.com/watch?v=e8Yw4aIG16Q>

Topics Covered in Today's Training

- 01 Why Time Series Analysis?
- 02 What is Time Series?
- 03 Components of Time Series
- 04 When not to use Time Series?
- 05 What is Stationarity?
- 06 ARIMA model
- 07 Demo: Forecast future



Why Time Series Analysis?

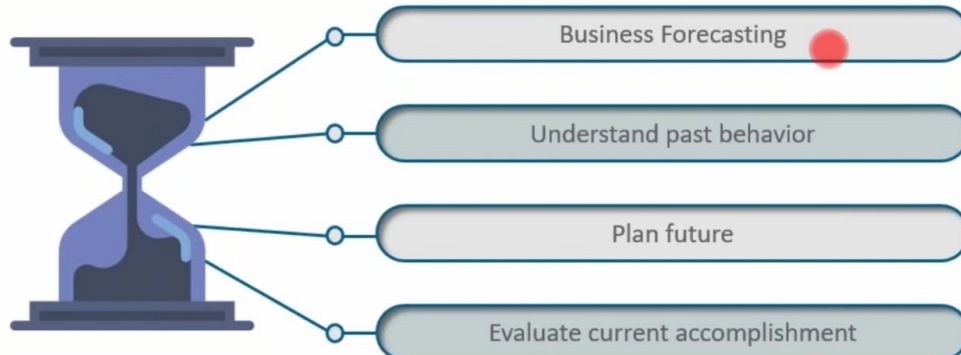
In this analysis, you just have one variable – **TIME**

You can analyse this **time series** data in order to extract meaningful statistics and other characteristics

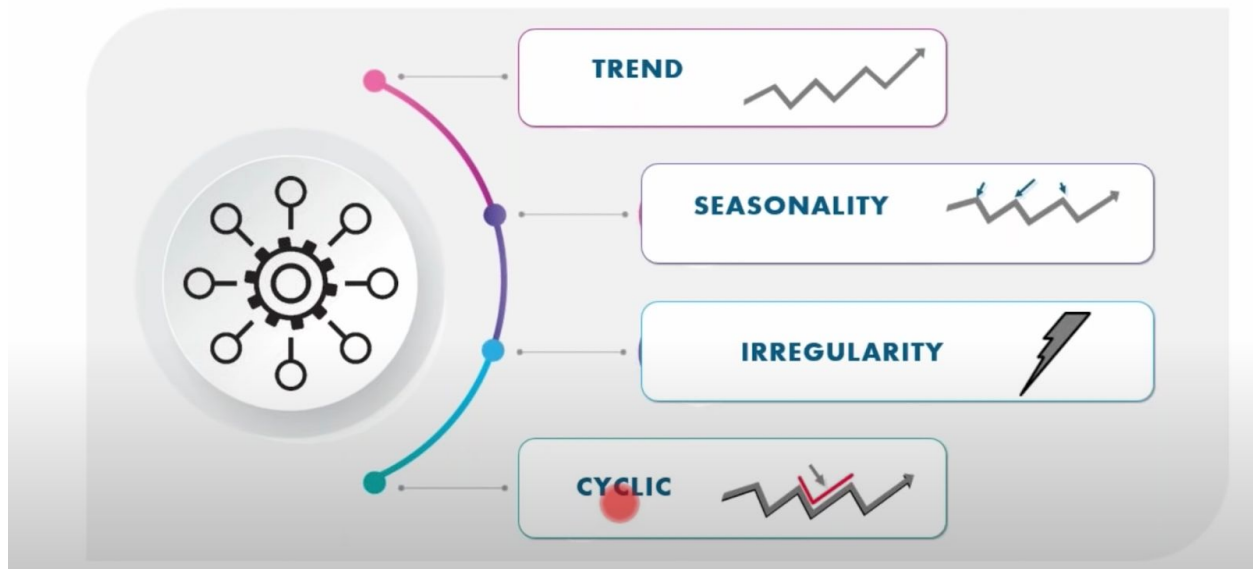


What Is Time Series?

- A time series is a set of observation taken at specified **times** usually at equal intervals
- It is used to **predict** the future values based on the **previous** observed values



Components Of Time Series



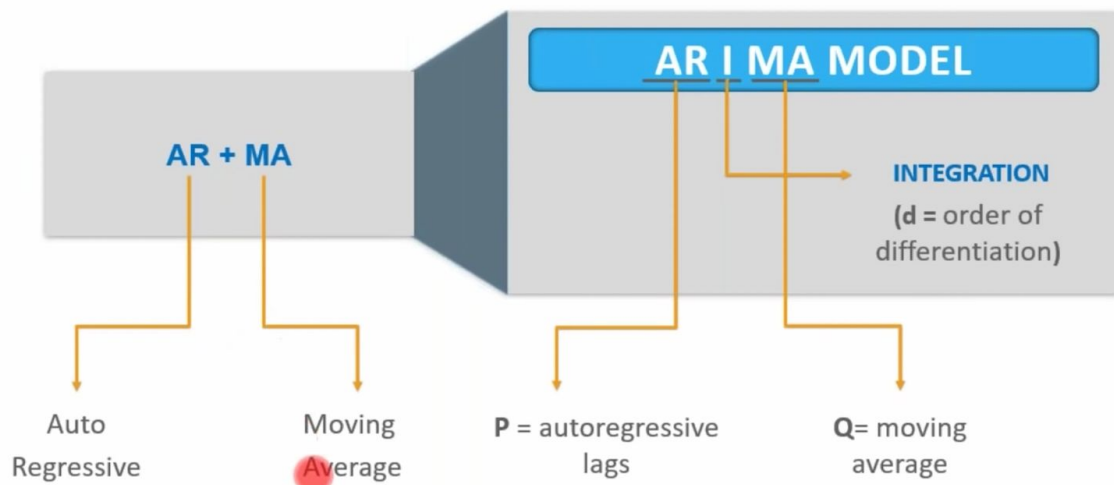
What Is Stationarity?

TS has a particular behaviour over time, there is a very high probability that it will follow the same in the **future**.

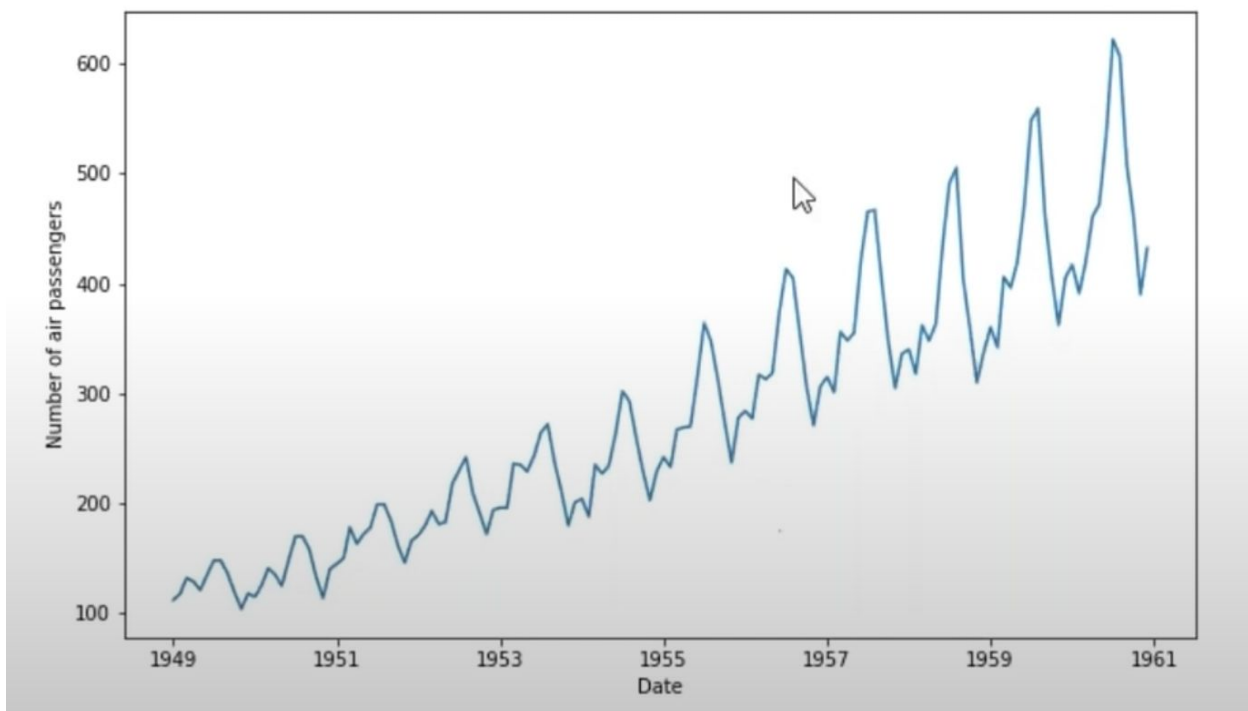
How to remove
Stationarity?

- 1 Constant mean
- 2 Constant Variance
- 3 Autocovariance that does not depend on time

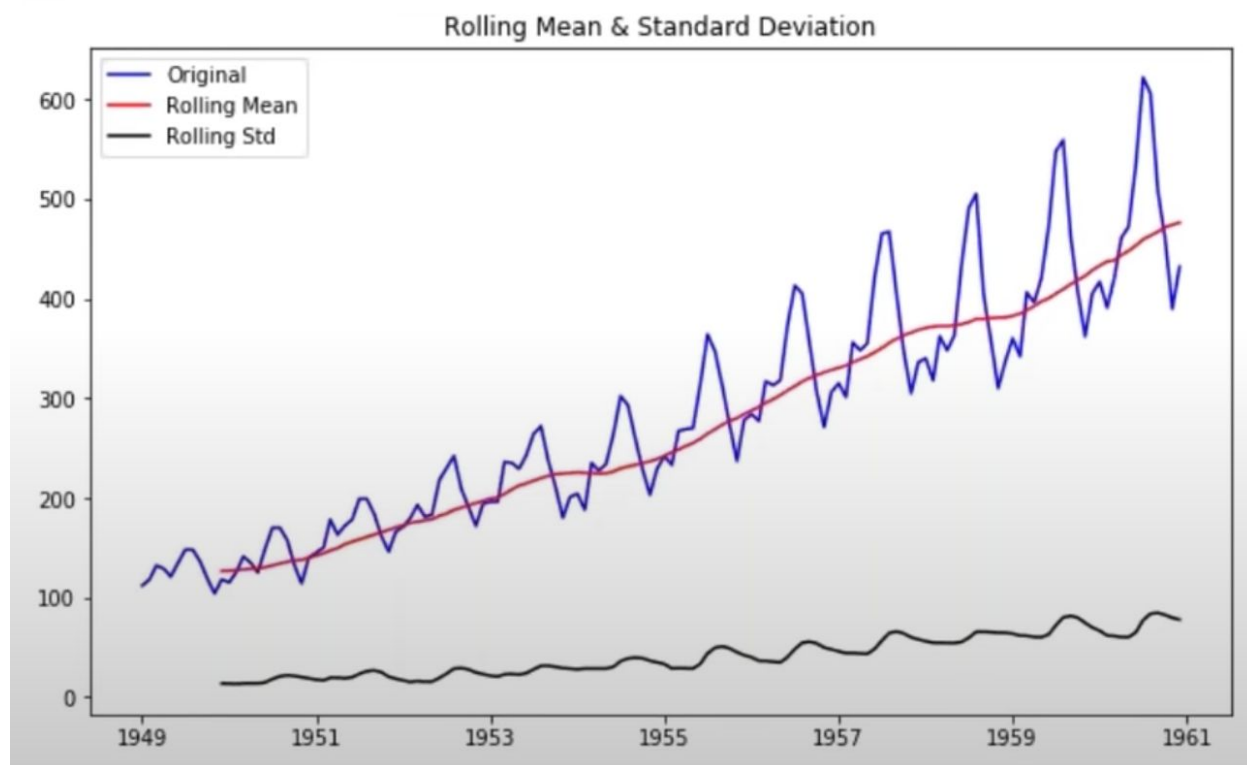
ARIMA MODEL



Not stationary data, conclusion reached by the non constant average (compare 1949 average with 1960 average). Graph below.

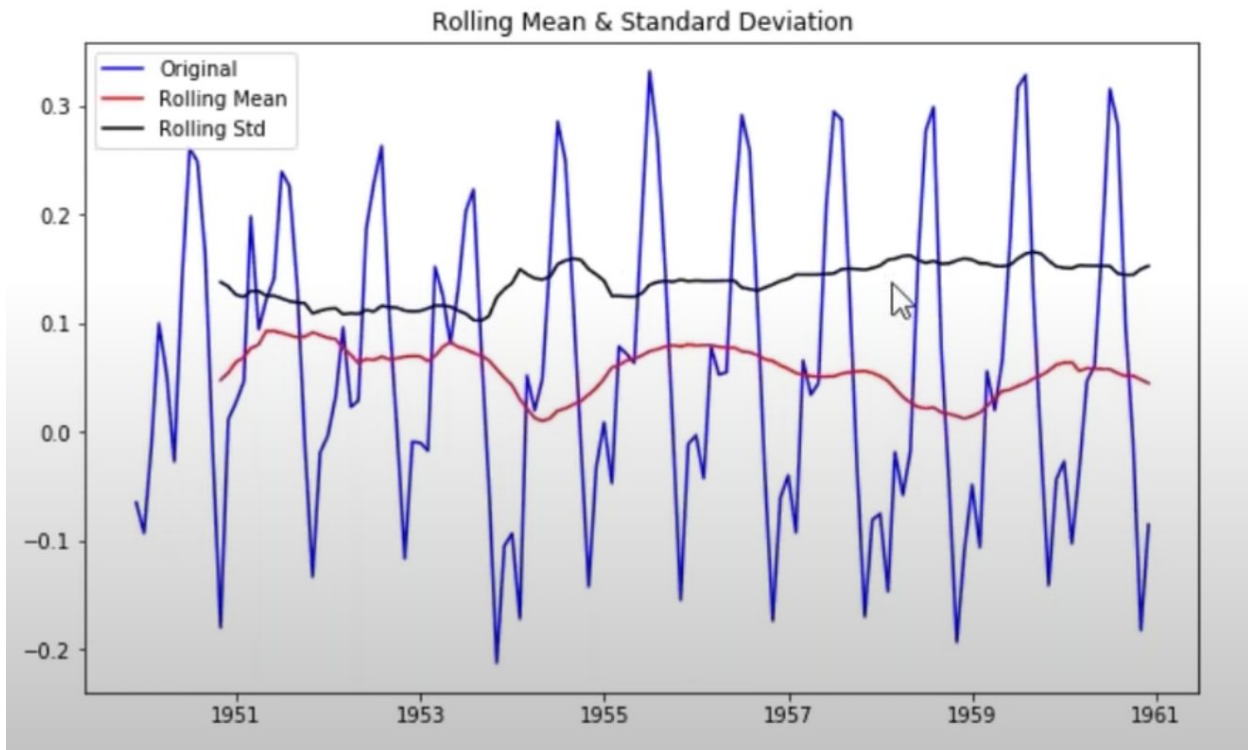


To be even more explicit, use the rolling average and standard deviation metrics.



```
Results of Dickey-Fuller Test:
Test Statistic           0.815369
p-value                   0.991880
#Lags Used                13.000000
Number of Observations Used 130.000000
Critical Value (1%)       -3.481682
Critical Value (5%)       -2.884042
Critical Value (10%)      -2.578770
dtype: float64
```

After transformation (log and subtracting the moving average over the log), we have:



Results of Dickey-Fuller Test:

Test Statistic	-3.162908
p-value	0.022135
#Lags Used	13.000000
Number of Observations Used	119.000000
Critical Value (1%)	-3.486535
Critical Value (5%)	-2.886151
Critical Value (10%)	-2.579896
dtype:	float64

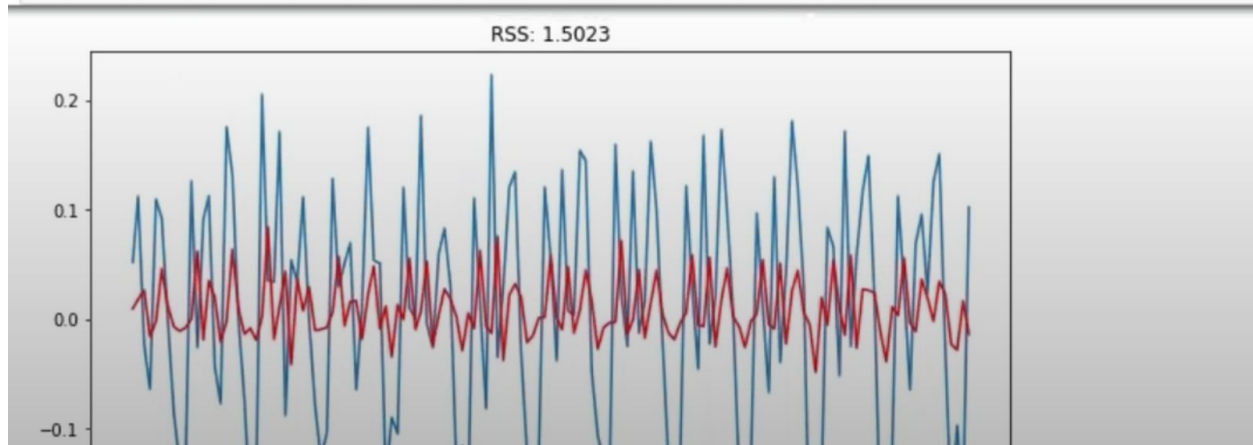
Training of the ARIMA model

```

from statsmodels.tsa.arima_model import ARIMA

#AR MODEL
model = ARIMA(indexedDataset_logScale, order=(0, 1, 0))
results_AR = model.fit(dis=-1)
plt.plot(datasetLogDiffShifting)
plt.plot(results_AR.fittedvalues, color='red')
plt.title('RSS: %.4f'% sum((results_AR.fittedvalues-datasetLogDiffShifting["#Passengers"])**2))
print('Plotting AR model')

```

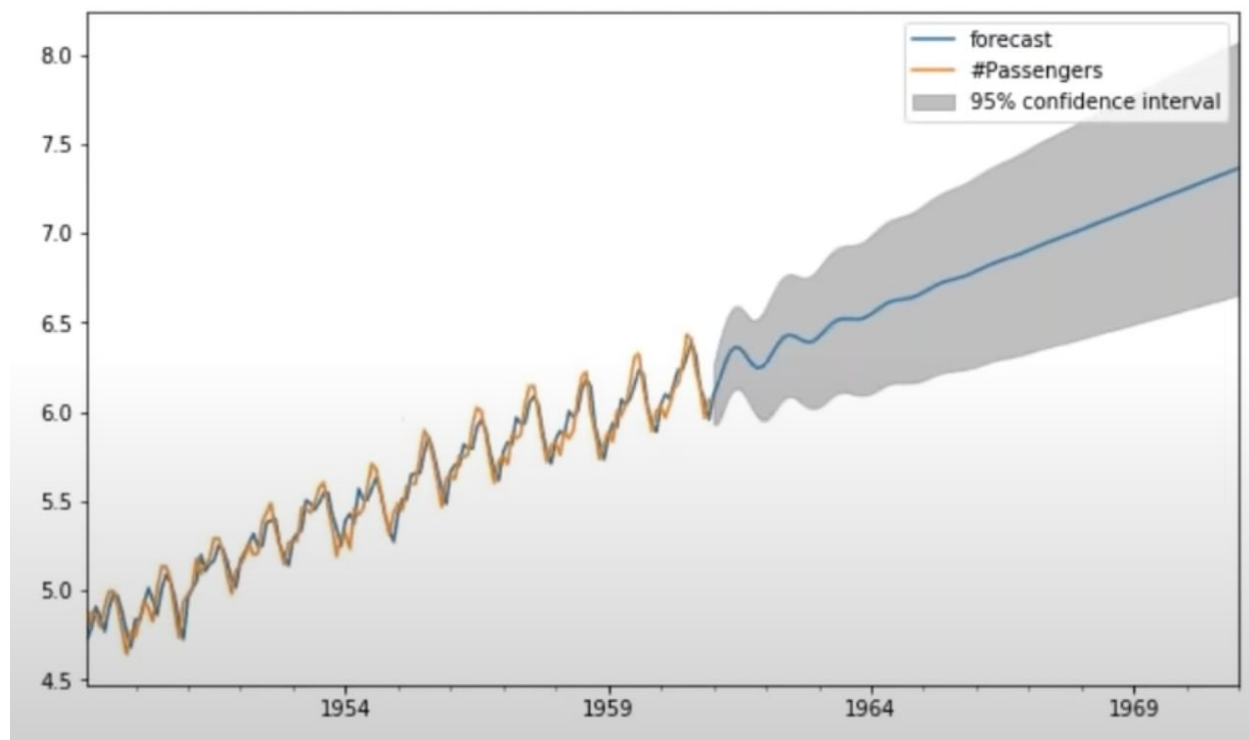


Then prediction for the next 10 years (120 months).

```

results_ARIMA.plot_predict(1, 264)
x=results_ARIMA.forecast(steps=120)

```



Introduction to Forecasting in Machine Learning and Deep Learning

InfoQ channel

<https://www.youtube.com/watch?v=bn8rVBulcFg>

Forecasting Methodologies

Classical, statistical (R: *forecast*)

- Autoregressive Integrated Moving Average (ARIMA)
- Exponential Smoothing Methods, e.g. Holt-Winters, Theta
- ...

Machine Learning

- Quantile Regression Forest (QRF)
- Support Vector Regression (SVR)
- Recurrent Neural Networks (RNNs)
- ...

Best model depends on, e.g.

- Amount of historical data
- Correlation with explanatory variables
- How the forecast will be used, e.g. interpretability needs, computational complexity

Compare multiple approaches.

UBER

Evaluation: Backtesting

Chronological testing

Sliding Window

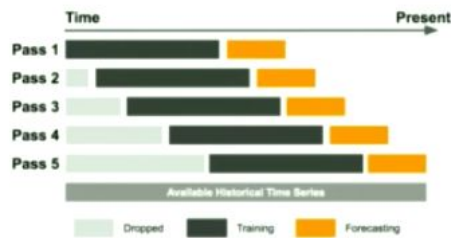


Figure 1: In the sliding window backtesting model, a fixed-size training window (in black) slides over the entire history of a time series and is repeatedly tested against a forecasting window (in orange) with older data points dropped.

Expanding Window

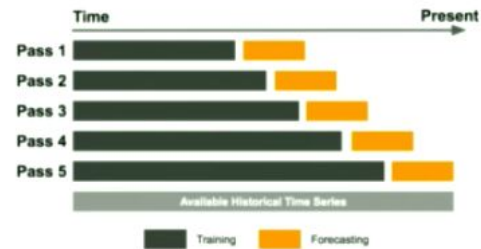
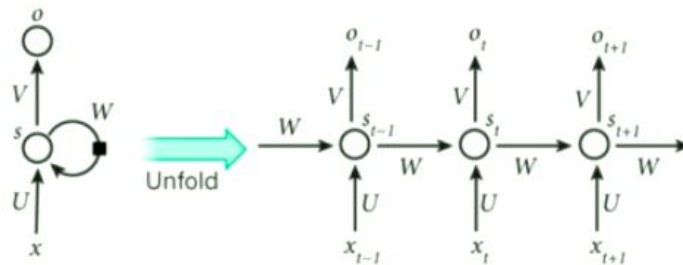


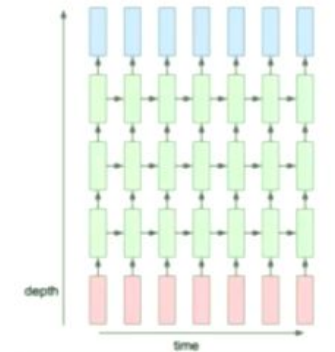
Figure 2: In the expanding window backtesting model, a training window (in black) expands over the entire history of a time series and is repeatedly tested against a forecasting window (in orange) without dropping older data points.

Evaluation Metric: Compare to Naive Forecast

Recurrent Neural Networks (RNNs)



*copyright <http://colah.github.io/>



RNNs can be multi-layer

Advantages

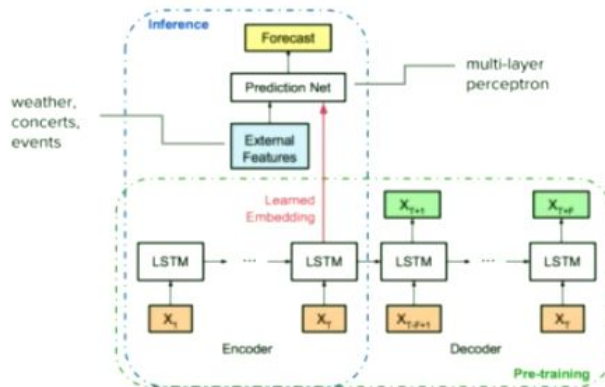
- end-to-end modeling,
- ease of incorporating exogenous variables,
- captures non-linear feature interactions
- automatic feature extraction.

Challenges

- need sufficient data for training,
- interpretability,
- need uncertainty estimates.

Forecasting Time Series for Extreme Events

LSTM Encoder-Decoder + Prediction Network



- **Encoder-Decoder:** two-layer LSTM cells, with 128 and 32 hidden states, respectively.
- **Prediction network:** three fully connected layers with tanh activation, with 128, 64, and 16 hidden units, respectively.
- **Raw data:** log-transformed to alleviate exponential effects.